


I/O-Algorithms

Lars Arge

Fall 2015

September 29, 2015

I/O-Model



- Parameters

N = # elements in problem instance

B = # elements that fits in disk block

M = # elements that fits in main memory

T = # output size in searching problem

- We often assume that $M > B^2$


- I/O**: Movement of block between memory and disk

Fundamental Bounds

	Internal	External
• Scanning:	N	$\frac{N}{B}$
• Sorting:	$N \log N$	$\frac{N}{B} \log_{M/B} \frac{N}{B}$
• Permuting	N	$\min\{N, \frac{N}{B} \log_{M/B} \frac{N}{B}\}$
• Searching:	$\log_2 N$	$\log_B N$


Fundamental Data Structures

- **B-trees**: Node degree $\Theta(B) \Rightarrow$ queries in $O(\log_B N + T/B)$
 - Rebalancing using split/fuse \Rightarrow updates in $O(\log_B N)$
- **Weight-balanced B-tress**: Weight rather than degree constraint
 $\Rightarrow \Omega(w(v))$ updates below v between rebalancing operations on v
- **Persistent B-trees**:
 - Update in current version in $O(\log_B N)$
 - Search in all previous versions in $O(\log_B N + T/B)$
- **Buffer trees**
 - Batching of operations to obtain $O(\frac{1}{B} \log_{M/B} \frac{N}{B})$ bounds
 $\Rightarrow O(\frac{N}{B} \log_{M/B} \frac{N}{B})$ construction algorithms



Interval Management


- **Problem:**
 - Maintain N intervals with **unique endpoints** dynamically such that stabbing query with point x can be answered efficiently



- As in (one-dimensional) B-tree case we are interested in
 - $O(N/B)$ space
 - $O(\log_B N)$ update
 - $O(\log_B N + T/B)$ query


Interval Management: Static Solution

- **Sweep** from left to right maintaining persistent B-tree
 - Insert interval when left endpoint is reached
 - Delete interval when right endpoint is reached



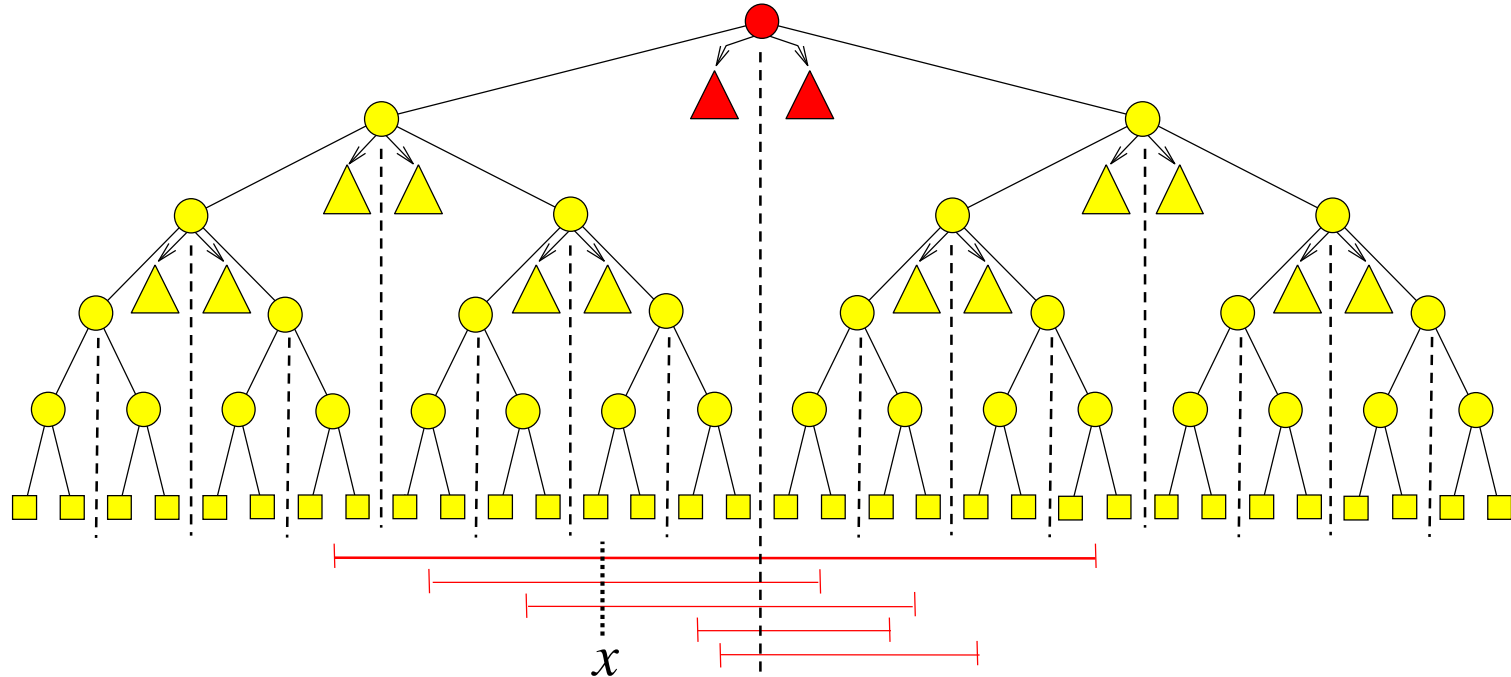
- Query x answered by reporting all intervals in B-tree at “time” x
 - $O(N/B)$ space
 - $O(\log_B N + T/B)$ query
 - $O(\frac{N}{B} \log_{M/B} \frac{N}{B})$ construction using buffer technique

Internal Interval Tree




- Base tree on endpoints – “slab” X_v associated with each node v
 - Interval stored in highest node v where it contains midpoint of X_v
 - Intervals I_v associated with v stored in
 - Left slab list sorted by left endpoint (search tree)
 - Right slab list sorted by right endpoint (search tree)
- ⇒ Linear space and $O(\log N)$ update (assuming fixed endpoint set)

Internal Interval Tree




- **Query** with x on left side of midpoint of X_{root}
 - Search **left slab list** left-right until finding non-stabbed interval
 - Recurse in left child
- $\Rightarrow O(\log N + T)$ query bound

Externalizing Interval Tree



- **Natural idea:**
 - Block tree
 - Use B-tree for **slab lists**
- Number of stabbed intervals in large slab list may be small (or zero)
 - We can be forced to do I/O in each of $O(\log N)$ nodes



Externalizing Interval Tree



- **Idea:**
 - Decrease fan-out to $\Theta(\sqrt{B}) \Rightarrow$ height remains $O(\log_B N)$
 - $\Theta(\sqrt{B})$ slabs define $\Theta(B)$ **multislabs**
 - Interval stored in two slab lists (as before) and one **multislab list**
 - Intervals in small multislab lists collected in **underflow structure**
 - Query answered in v by looking at 2 slab lists and not $O(\log B)$


External Interval Tree

- Base tree: Weight-balanced B-tree with branching parameter $\frac{1}{4}\sqrt{B}$ and leaf parameter B on endpoints
 - Interval stored in highest node v where it contains slab boundary
- Each internal node v contains:
 - Left slab list for each of $\Theta(\sqrt{B})$ slabs
 - Right slab lists for each of $\Theta(\sqrt{B})$ slabs
 - $\Theta(B)$ multislab lists
 - Underflow structure
- Interval in set I_v of intervals associated with v stored in
 - Left slab list of slab containing left endpoint
 - Right slab list of slab containing right endpoint
 - Widest multislab list it spans
- If $< B$ intervals in multislab list they are instead stored in underflow structure (\Rightarrow contains $\leq B^2$ intervals)




External Interval tree

- Each leaf contains $< B/2$ intervals (unique endpoint assumption)
 - Stored in one block
 - **Slab lists** implemented using B-trees
 - $O(1 + T_v/B)$ query
 - Linear space
 - * We may “wasted” a block for each of the $\Theta(\sqrt{B})$ lists in node
 - * But only $\Theta(\frac{N}{B\sqrt{B}})$ internal nodes
 - **Underflow structure** implemented using static structure
 - $O(\log_B B^2 + T_v/B) = O(1 + T_v/B)$ query
 - Linear space
- ⇓
- Linear space




External Interval Tree

- Query with x
 - Search down tree for x while in node v reporting all intervals in I_v stabbed by x



- In node v
 - Query two slab lists
 - Report all intervals in relevant multislab lists
 - Query underflow structure




- Analysis:
 - Visit $O(\log_B N)$ nodes
 - Query slab lists
 - Query multislab lists
 - Query underflow structure
- $$\left. \begin{array}{l} \text{– Visit } O(\log_B N) \text{ nodes} \\ \text{– Query } \text{slab lists} \\ \text{– Query } \text{multislab lists} \\ \text{– Query } \text{underflow structure} \end{array} \right\} O(1 + T_v/B) \Rightarrow O(\log_B N + T/B)$$

External Interval Tree

- **Update** – ignoring base tree update/rebalancing:

- Search for relevant node
 - Update two **slab lists**
 - Update **multislabs list** or **underflow structure**
- $\left. \begin{array}{l} \text{Search for relevant node} \\ \text{Update two slab lists} \end{array} \right\} O(\log_B N)$



- Update of **underflow structure** in $O(I)$ I/Os amortized:
 - Maintain update block with $\leq B$ updates
 - Check of update block adds $O(I)$ I/Os to query bound
 - Rebuild structure when B updates have been collected using $O(\frac{B^2}{B} \log_B B^2) = O(B)$ I/Os (**Global rebuilding**)



Update in $O(\log_B N)$ I/Os amortized

External Interval Tree

- **Note:**
 - Insert may increase number of intervals in **underflow structure** for some **multislab** to B
 - Delete may decrease number of intervals in **multislab** to B

⇓

Need to move B intervals to/from **multislab/underflow structure**
- We only move
 - Intervals from **multislab list** when decreasing to size $B/2$
 - Intervals to **multislab list** when increasing to size B

⇓

$O(I)$ I/Os amortized used to move intervals

Base Tree Update

- Before **inserting** new interval we insert new endpoints in base tree using $O(\log_B N)$ I/Os


– Leads to rebalancing using splits



Boundary in v becomes boundary in $parent(v)$





Intervals need to be moved




- Move intervals (update secondary structures) in $O(w(v))$ I/Os
 $\Rightarrow O(1)$ amortized split bound (weight balanced B-tree)
 $\Rightarrow O(\log_B N)$ amortized insert bound

Splitting Interval Tree Node

- When v splits we may need to move $O(w(v))$ intervals
 - Intervals in v containing boundary
 - Intervals in $parent(v)$ with endpoints in X_v containing boundary
- Intervals move to two new **slab** and **multislabs** lists in $parent(v)$




Splitting Interval Tree Node



- Moving intervals in v in $O(w(v))$ I/Os
 - Collect in left order (and remove) by scanning left **slab lists**
 - Collect in right order (and remove) by scanning right **slab lists**
 - Remove **multislabs lists** containing boundary
 - Remove from **underflow structure** by rebuilding it
 - Construct lists and **underflow structure** for v' and v'' similarly


Splitting Interval Tree Node



- Moving intervals in $parent(v)$ in $O(w(v))$ I/Os
 - Collect in left order by scanning left **slab list**
 - Collect in right order by scanning right **slab list**
 - Merge with intervals collected in $v \Rightarrow$ two new **slab lists**
 - Construct new **multislab lists** by splitting relevant **multislab list**
 - Insert intervals in small **multislab lists** in **underflow structure**

External Interval Tree


- Split in $O(I)$ I/Os amortized
 - Space: $O(N/B)$
 - Query: $O(\log_B N + T/B)$
 - Insert: $O(\log_B N)$ I/Os amortized



- **Deletes** in $O(\log_B N)$ I/Os amortized using **global rebuilding**:
 - Delete interval as previously using $O(\log_B N)$ I/Os
 - Mark relevant endpoint as deleted
 - Rebuild structure in $O(N \log_B N)$ after $N/2$ deletes
- Note: Deletes can also be handled using **fuse** operations

External Interval Tree

- External interval tree
 - Space: $O(N/B)$
 - Query: $O(\log_B N + T/B)$
 - Updates: $O(\log_B N)$ I/Os amortized



- Removing amortization:
 - Moving intervals to/from **underflow structure**
 - Delete global rebuilding
 - Underflow structure update
 - Base node tree splits


Perform operations/construction lazily

Move lazily – complicated:

- Interference
- Queries


Summary/Conclusion: Interval Management

- Interval management corresponds to simple form of $2d$ range search
 - Diagonal corner queries
- We obtained the same bounds as for the $1d$ case
 - Space: $O(N/B)$
 - Query: $O(\log_B N + T/B)$
 - Updates: $O(\log_B N)$ I/Os




Summary/Conclusion: Interval Management

- Main problem in designing structure:
 - Binary \rightarrow large fan-out
- Large fan-out resulted in the need for
 - Multislabs and multislabs lists
 - Underflow structure to avoid $O(B)$ -cost in each node
- General solution techniques:
 - **Filtering**: Charge part of query cost to output
 - **Bootstrapping**:
 - * Use $O(B^2)$ size structure in each internal node
 - * Constructed using persistence
 - * Dynamic using global rebuilding
 - **Weight-balanced B-tree**: Split/fuse in amortized $O(1)$




Three-Sided Range Queries

- Interval management: “1.5 dimensional” search




- More general 2d problem: **Dynamic 3-sided range searching**
 - Maintain set of points in plane such that given query (q_1, q_2, q_3) , all points (x, y) with $q_1 \leq x \leq q_2$ and $y \geq q_3$ can be found efficiently




Three-Sided Range Queries

- Report all points (x,y) with $q_1 \leq x \leq q_2$ and $y \geq q_3$
- **Static solution:**
 - Sweep top-down inserting x in persistent B-tree at (x,y)
 - Answer query by performing range query with $[q_1, q_2]$ in B-tree at q_3
- **Optimal:**
 - $O(N/B)$ space
 - $O(\log_B N + T/B)$ query
 - $O(\frac{N}{B} \log_{M/B} \frac{N}{B})$ construction
- **Dynamic?** ... in internal memory: **priority search tree**




Internal Priority Search Tree



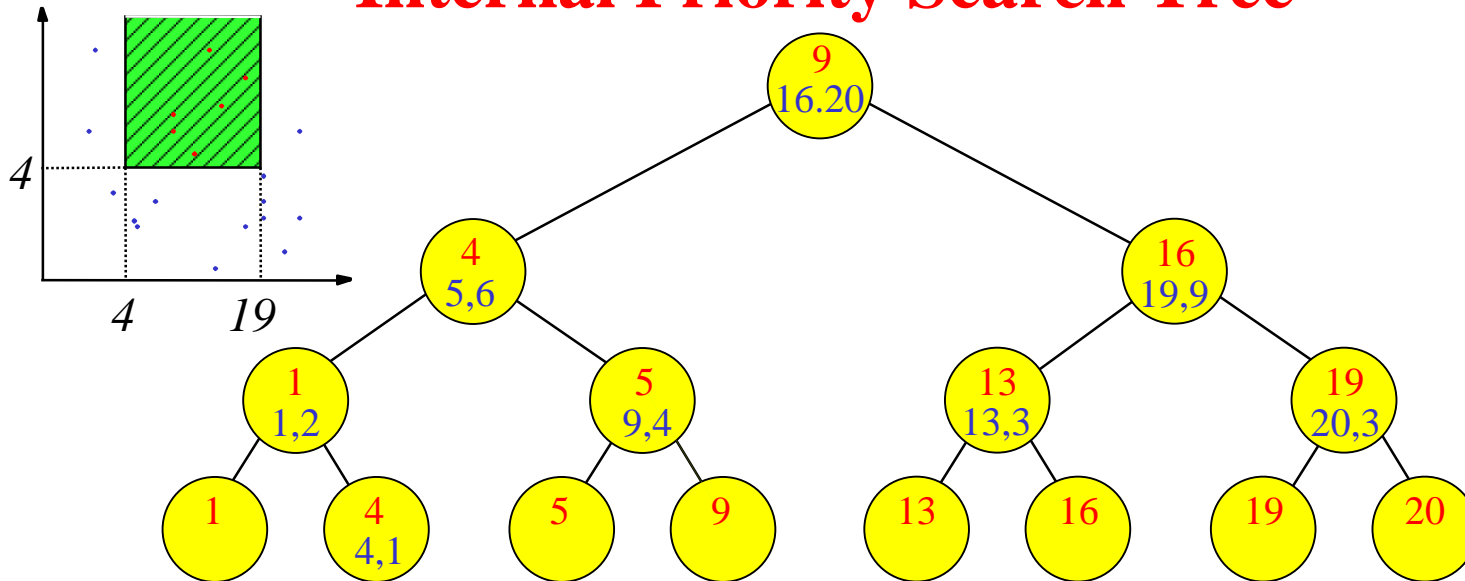
- **Base tree on x -coordinates** with nodes augmented with points
- **Heap on y -coordinates**
 - Decreasing y values on root-leaf path
 - (x,y) on path from root to leaf holding x
 - If v holds point then $parent(v)$ holds point

Internal Priority Search Tree



- Linear space
 - **Insert** of (x,y) (assuming fixed x -coordinate set):
 - Compare y with y -coordinate in root
 - Smaller: Recursively insert (x,y) in subtree on path to x
 - Bigger: Insert in root and recursively insert old point in subtree
- $\Rightarrow O(\log N)$ update


Internal Priority Search Tree



- **Query** with (q_1, q_2, q_3) starting at root v :
 - Report point in v if satisfying query
 - Visit both children of v if point reported
 - Always visit child(s) of v on path(s) to q_1 and q_2


$\Rightarrow O(\log N + T)$ query

Externalizing Priority Search Tree



- **Natural idea:** Block tree
 - **Problem:**
 - $O(\log_B N)$ I/Os to follow paths to q_1 and q_2
 - But $O(T)$ I/Os may be used to visit other nodes (“overshooting”)
- $\Rightarrow O(\log_B N + T)$ query

Externalizing Priority Search Tree




- **Solution idea:**
 - Store B points in each node \Rightarrow
 - * $O(B^2)$ points stored in each supernode
 - * B output points can pay for “overshooting”
 - **Bootstrapping:**
 - * Store $O(B^2)$ points in each supernode in static structure

External Priority Search Tree

- **Base tree**: Weight-balanced B-tree with branching parameter $B/4$ and leaf parameter B on x -coordinates
- Points in “**heap order**”:
 - Root stores B top points for each of the $\Theta(B)$ child slabs
 - Remaining points stored recursively
- Points in each node stored in “ **B^2 -structure**”
 - Persistent B-tree structure for static problem




Linear space



External Priority Search Tree

- **Query** with (q_1, q_2, q_3) starting at root v :
 - Query B^2 -structure and report points satisfying query
 - Visit child v if
 - * v on path to q_1 or q_2
 - * All points corresponding to v satisfy query




External Priority Search Tree

- **Analysis:**

- $O(\log_B B^2 + T_v/B) = O(1 + T_v/B)$ I/Os used to visit node v
- $O(\log_B N)$ nodes on path to q_1 or q_2
- For each node v not on path to q_1 or q_2 visited, B points reported in $parent(v)$


⇓

$O(\log_B N + T/B)$ query



External Priority Search Tree

- **Insert** (x,y) (ignoring insert in base tree - rebalancing):
 - Find relevant node u :
 - * Query B^2 -structure to find B points in root corresponding to node u on path to x
 - * If y smaller than y -coordinates of all B points then recursively search in u
 - Insert (x,y) in B^2 -structure of v
 - If B^2 -structure contains $>B$ points for child u , remove lowest point and insert recursively in u
- **Delete**: Similarly




External Priority Search Tree

- **Analysis:**

- Update visits $O(\log_B N)$ nodes
- B^2 -structure queried/updated in each node
 - * One query
 - * One insert and one delete

- **B^2 -structure analysis:**

- Query: $O(\log_B B^2 + B/B) = O(1)$
- Update: $O(1)$ using **global rebuilding**
 - * Store updates in update block
 - * Rebuild after B updates using $O(\frac{B^2}{B} \log_{M/B} \frac{B^2}{B}) = O(B)$ I/Os



$O(\log_B N)$ I/O updates

Dynamic Base Tree

- **Deletion:**


- Delete point as previously
- Delete x -coordinate from base tree using **global rebuilding**

$\Rightarrow O(\log_B N)$ I/Os amortized

- **Insertion:**


- Insert x -coordinate in base tree and rebalance (using **splits**)
- Insert point as previously

- **Split:** Boundary in v becomes boundary in $parent(v)$



Dynamic Base Tree

- **Split**: When v splits B new points needed in $\text{parent}(v)$
- One point obtained from v' (v'') using “**bubble-up**” operation:
 - Find top point p in v'
 - Insert p in B^2 -structure
 - Remove p from B^2 -structure of v'
 - Recursively bubble-up point to v'
- **Bubble-up** in $O(\log_B w(v))$ I/Os
 - Follow one path from v to leaf
 - Uses $O(1)$ I/O in each node




Split in $O(B \log_B w(v)) = O(w(v))$ I/Os

Dynamic Base Tree

- $O(I)$ amortized split cost:
 - Cost: $O(w(v))$
 - Weight balanced base tree: $\Omega(w(v))$ inserts below v between splits



- **External Priority Search Tree**
 - Space: $O(N/B)$
 - Query: $O(\log_B N + T/B)$
 - Updates: $O(\log_B N)$ I/Os amortized





- Amortization can be removed from update bound in several ways
 - Utilizing lazy rebuilding


Summary/Conclusion: Priority Search Tree

- We have now discussed structures for **special cases** of two-dimensional range searching

- Space: $O(N/B)$
- Query: $O(\log_B N + T/B)$
- Updates: $O(\log_B N)$



- Cannot be obtained for general (4-sided) **2d range searching**:
 - $O(\log_B^c N)$ query requires $\Omega(\frac{N}{B} \frac{\log_B N}{\log_B \log_B N})$ space
 - $O(\frac{N}{B})$ space requires $\Omega(\sqrt{N/B})$ query



References

- **External Memory Geometric Data Structures**

Lecture notes by Lars Arge.

– Section 6-7